

Introduction to the Message Passing Interface (MPI)

Collective Communication


Dr. Juraj Kardos, Dr. Pratyuksh Bansal, Prof. Olaf Schenk
(TA: Tim Holt, Malik Lechekhab)

Faculty of Informatics
USI Lugano, Switzerland
November 3, 2021

Chap. 5 Collective Communication

1. MPI Overview

- one program on several processors
work and data distribution




2. Process model and language bindings

- starting several MPI processes

```
MPI_Init()  
MPI_Comm_rank()
```

3. Messages and point-to-point communication

- the MPI processes can communicate




4. Non-blocking communication

- to avoid idle time and deadlocks



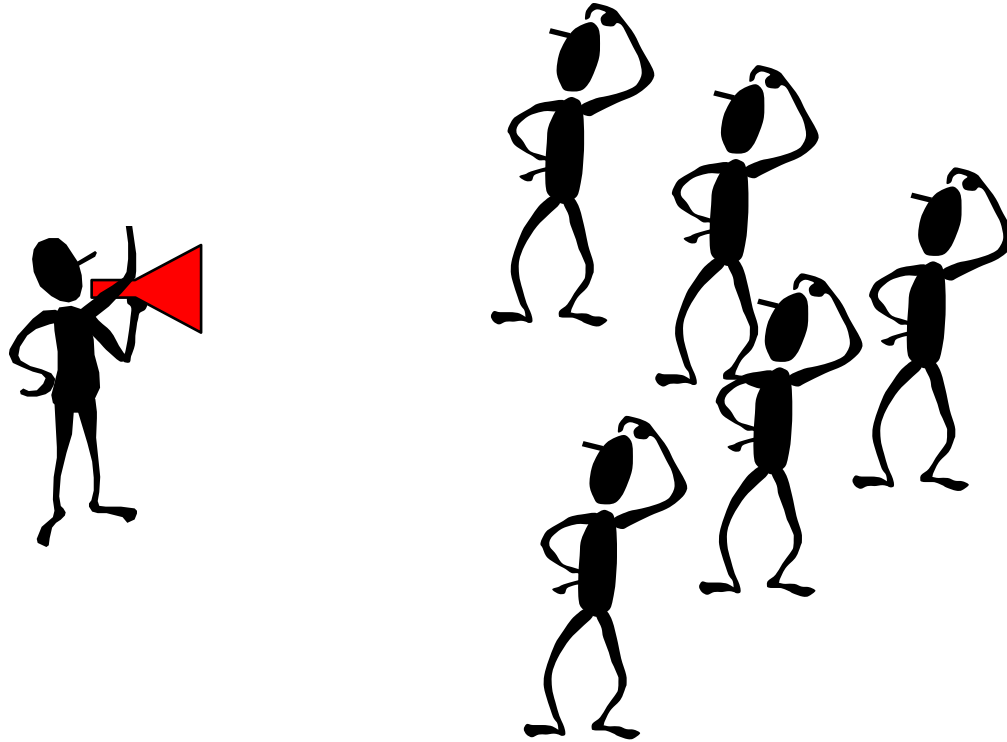
5. Collective communication

- e.g., broadcast



Broadcast

- A one-to-many communication.

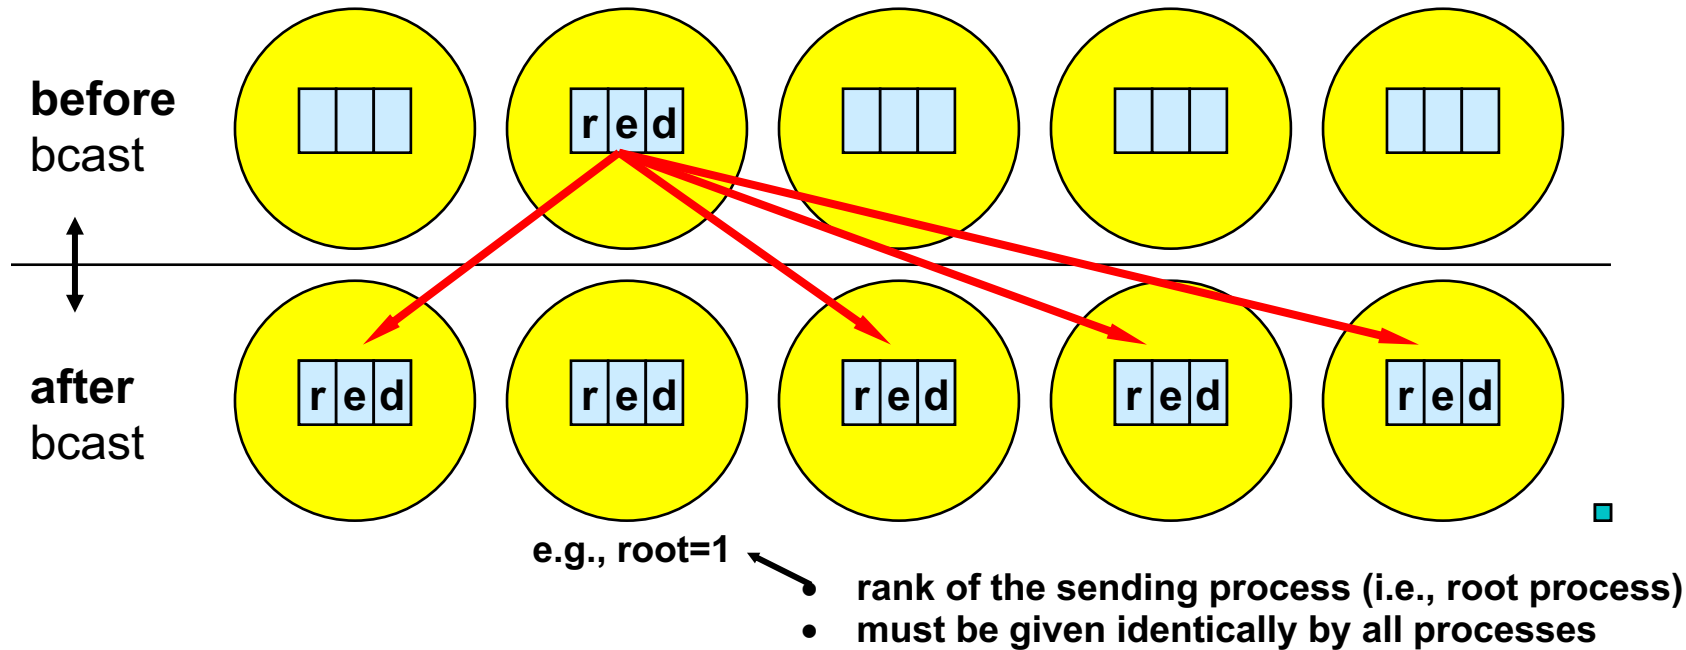


Collective Communication

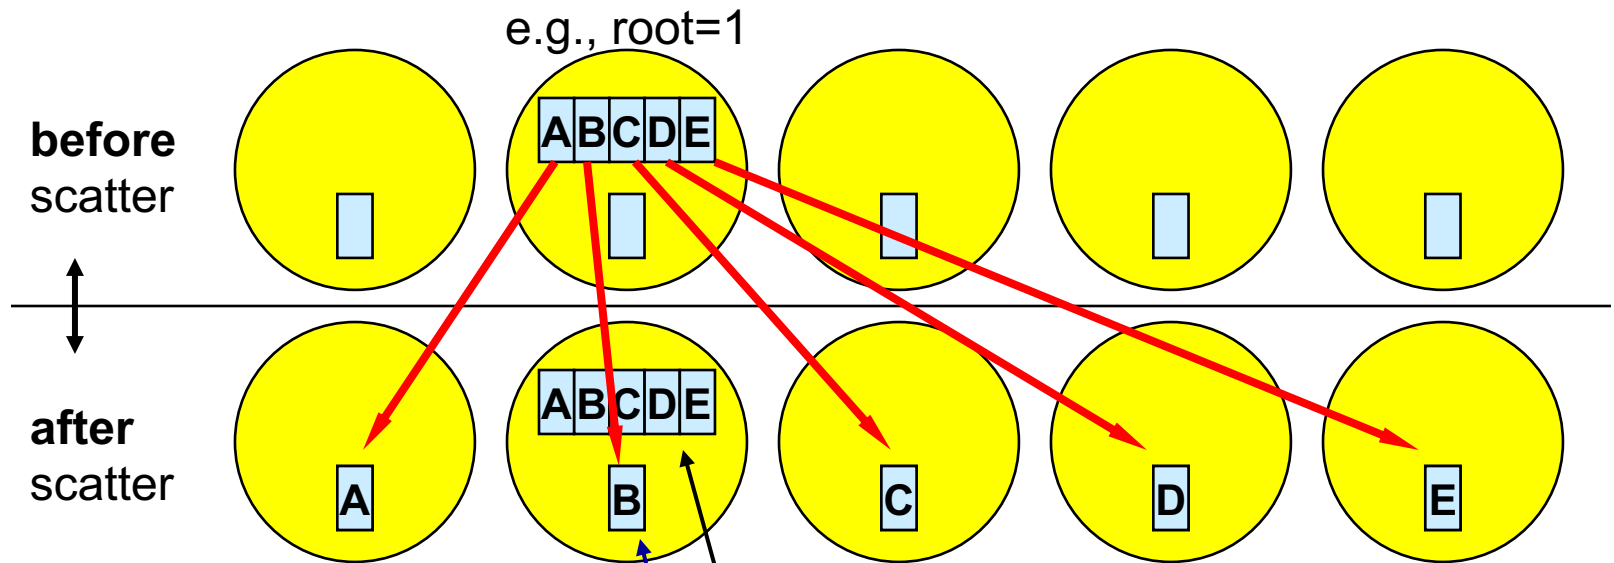
- Communications involving a group of processes.
- Called by all processes in a communicator.
- Examples:
 - Broadcast, scatter, gather.
 - Global sum, global maximum, etc.

Broadcast

- C: int **MPI_Bcast**(void *buf, int count, **MPI_Datatype** datatype, int root, **MPI_Comm** comm)



Scatter

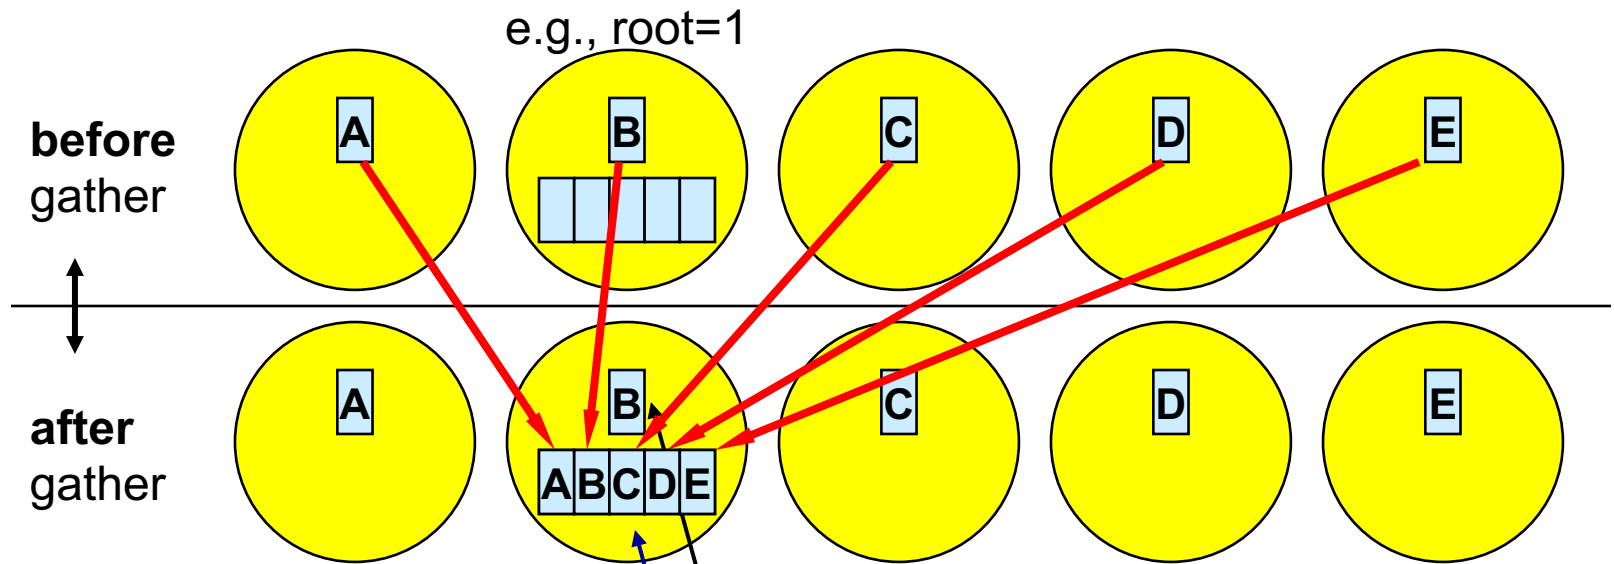


- C: `int MPI_Scatter(void *sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)`

Example:

`MPI_Scatter(sbuf, 1, MPI_CHAR, rbuf, 1, MPI_CHAR, 1, MPI_COMM_WORLD)`

Gather



- C: `int MPI_Gather(void *sendbuf, int sendcount, MPI_Datatype sendtype,
void *recvbuf, int recvcount, MPI_Datatype
recvtype,
int root, MPI_Comm comm)`

Global Reduction Operations

- To perform a global reduce operation across all members of a group.
- $d_0 \circ d_1 \circ d_2 \circ d_3 \circ \dots \circ d_{s-2} \circ d_{s-1}$
 - d_i = data in process rank i
 - single variable, or
 - vector
 - \circ = associative operation
 - Example:
 - global sum or product
 - global maximum or minimum
 - global user-defined operation
- floating point rounding may depend on usage of associative law:
 - $[(d_0 \circ d_1) \circ (d_2 \circ d_3)] \circ [\dots \circ (d_{s-2} \circ d_{s-1})]$
 - $(((((d_0 \circ d_1) \circ d_2) \circ d_3) \circ \dots) \circ d_{s-2}) \circ d_{s-1})$

Example of Global Reduction


- Global integer sum.
- Sum of all inbuf values should be returned in *resultbuf*.
- C: int root=0
 MPI_Reduce(&inbuf, &*resultbuf*, 1, **MPI_INT**, **MPI_SUM**,
 root, **MPI_COMM_WORLD**);
- The result is only placed in *resultbuf* at the root process.

Predefined Reduction Operation Handles

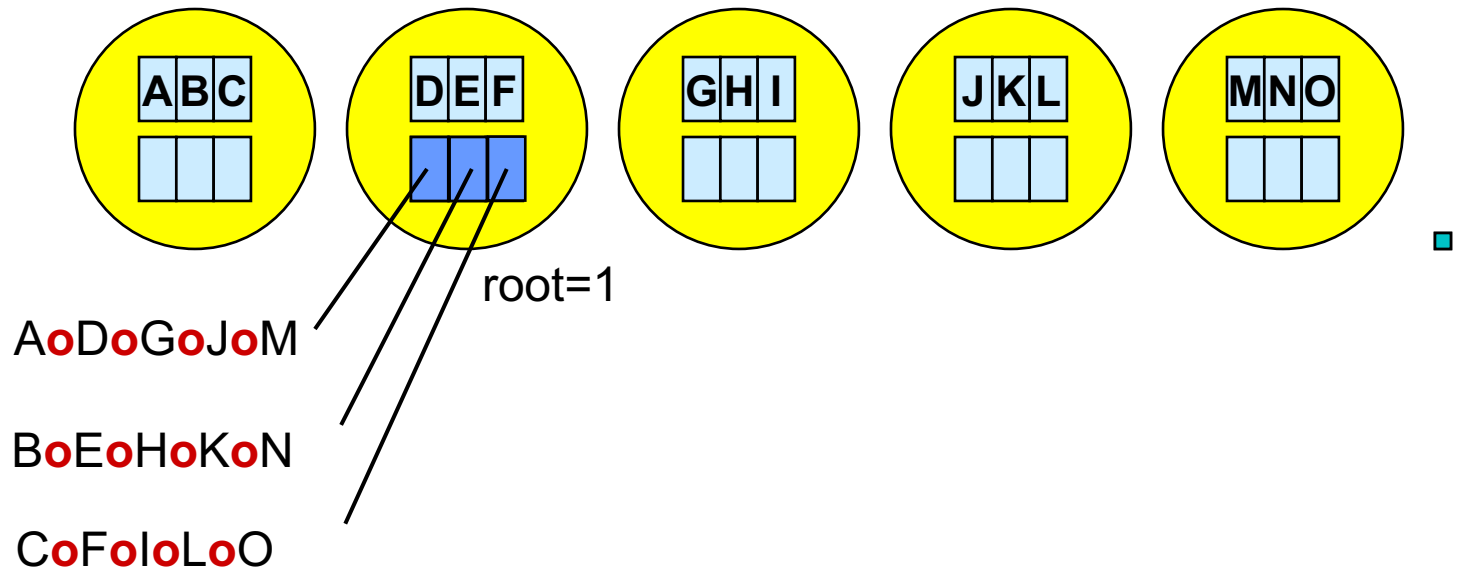
Predefined operation handle	Function
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI_LAND	Logical AND
MPI_BAND	Bitwise AND
MPI_LOR	Logical OR
MPI_BOR	Bitwise OR
MPI_LXOR	Logical exclusive OR
MPI_BXOR	Bitwise exclusive OR
MPI_MAXLOC	Maximum and location of the maximum
MPI_MINLOC	Minimum and location of the minimum

MPI_REDUCE

before **MPI_REDUCE**



after



Variants of Reduction Operations

- **MPI_ALLREDUCE**

- no root,
- returns the result in all processes

- **MPI_SCAN**

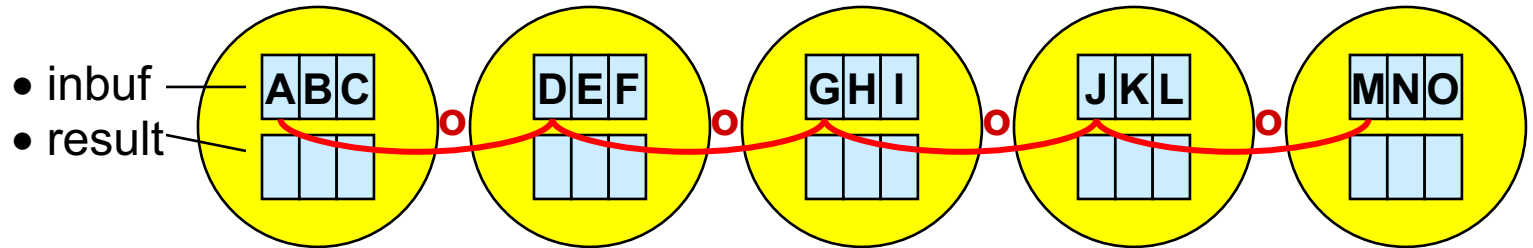
- prefix reduction
- result at process with rank i :=
reduction of inbuf-values from rank 0 to rank i

- **MPI_EXSCAN**

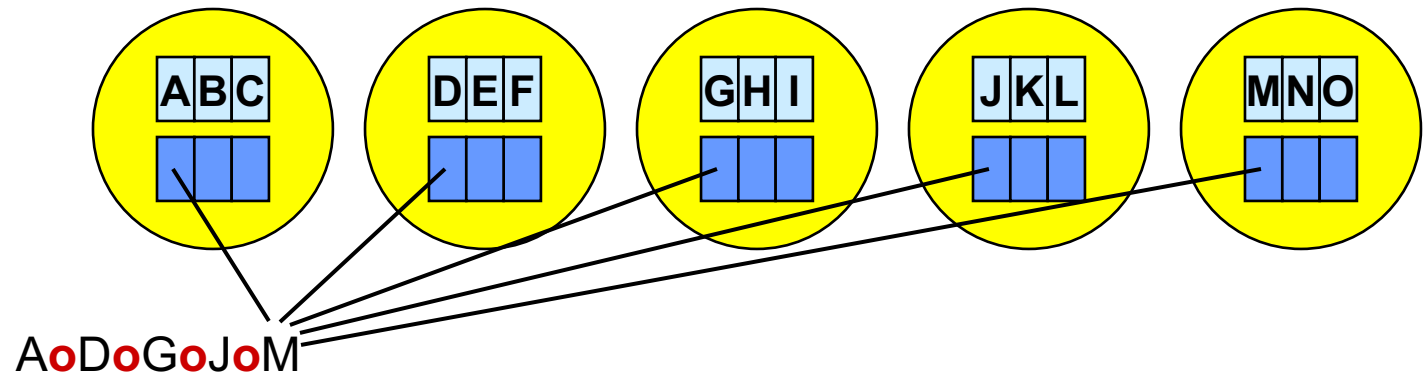
- result at process with rank i :=
reduction of inbuf-values from rank 0 to rank $i-1$

MPI_ALLREDUCE

before MPI_ALLREDUCE



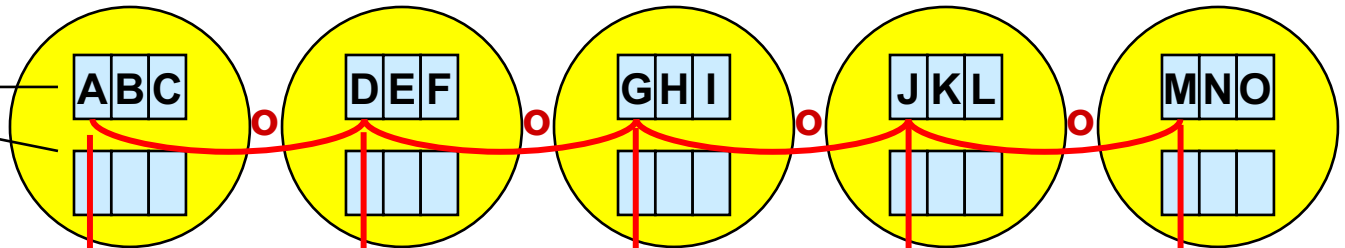
after



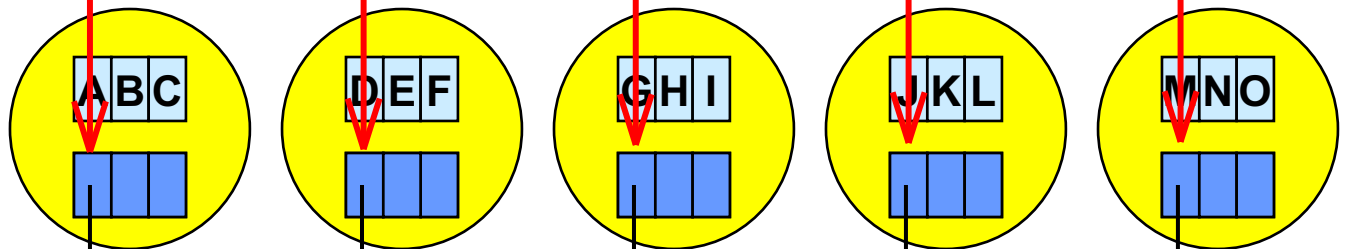
MPI_SCAN and MPI_EXSCAN

before the call

- inbuf —
- result —



after



MPI_SCAN: A AoD AoDoG AoDoGoJ AoDoGoJoM

MPI_EXSCAN: - A AoD AoDoG AoDoGoJ

done in parallel



In Class Exercise: PI

- You are given the code template, that calculates PI value in pi.c
- The PI value is calculated by numerically solving the integral of arctangent
 - Each process calculates approximate area of its part of integral
 - In the end, partial results shall be summed together
- Use MPI reduction call to collect the grand total integral sum on one rank
- Print out the resulting approximation of PI value

In Class Exercise: PI

- show solution pi.c